

## ✓ Project Name: Heart Attack Risk Predictor

In this project we will Make an app which will help us predict the risk of a Heart Attack a person have.

We will do use various Algorithms to predict the result and see which one suits best and then we will use Auto ML Library EVAL ML to predict the results.

We will do the following things:

- Data Analysis
- Feature Engineering
- Satandardization
- Model Building
- Predictions

## ✓ Let us import the necessary liabraries and read our DataSet

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

## ✓ Let us import our Data Set

```
df= pd.read_csv("/content/heart.csv")
```

```
df= df.drop(['oldpeak','slp','thall'],axis=1)
```

```
df.info()
```


```
↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
```

```

1  sex      303 non-null  int64
2  cp       303 non-null  int64
3  trtbps   303 non-null  int64
4  chol     303 non-null  int64
5  fbs      303 non-null  int64
6  restecg  303 non-null  int64
7  thalachh 303 non-null  int64
8  exng     303 non-null  int64
9  caa      303 non-null  int64
10 output   303 non-null  int64
dtypes: int64(11)
memory usage: 26.2 KB

```

```
df.head()
```



	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	caa	output
0	63	1	3	145	233	1	0	150	0	0	1
1	37	1	2	130	250	0	1	187	0	0	1
2	41	0	1	130	204	0	0	172	0	0	1
3	56	1	1	120	236	0	1	178	0	0	1
4	57	0	0	120	354	0	1	163	1	0	1

## ✓ Data Analysis

### ✓ Understanding our DataSet:

Age : Age of the patient

Sex : Sex of the patient

exang: exercise induced angina (1 = yes; 0 = no)

ca: number of major vessels (0-3)

cp : Chest Pain type chest pain type

- Value 0: typical angina
- Value 1: atypical angina
- Value 2: non-anginal pain
- Value 3: asymptomatic

trtbps : resting blood pressure (in mm Hg)

chol : cholestoral in mg/dl fetched via BMI sensor

fbs : (fasting blood sugar > 120 mg/dl) (1 = true; 0 = false)

rest\_ecg : resting electrocardiographic results

- Value 0: normal
- Value 1: having ST-T wave abnormality (T wave inversions and/or ST elevation or depression of > 0.05 mV)
- Value 2: showing probable or definite left ventricular hypertrophy by Estes' criteria

thalach : maximum heart rate achieved

target : 0= less chance of heart attack 1= more chance of heart attack

df.shape

↩ (303, 11)

df.isnull().sum()

↩

	0
age	0
sex	0
cp	0
trtbps	0
chol	0
fbs	0
restecg	0
thalachh	0
exng	0
caa	0
output	0

dtype: int64

✓ As we can see there are no null values in our Data Set

df.corr()



	age	sex	cp	trtbps	chol	fbs	restecg	thalachh
age	1.000000	-0.098447	-0.068653	0.279351	0.213678	0.121308	-0.116211	-0.398522
sex	-0.098447	1.000000	-0.049353	-0.056769	-0.197912	0.045032	-0.058196	-0.044020
cp	-0.068653	-0.049353	1.000000	0.047608	-0.076904	0.094444	0.044421	0.295762
trtbps	0.279351	-0.056769	0.047608	1.000000	0.123174	0.177531	-0.114103	-0.046698
chol	0.213678	-0.197912	-0.076904	0.123174	1.000000	0.013294	-0.151040	-0.009940
fbs	0.121308	0.045032	0.094444	0.177531	0.013294	1.000000	-0.084189	-0.008567
restecg	-0.116211	-0.058196	0.044421	-0.114103	-0.151040	-0.084189	1.000000	0.044123
thalachh	-0.398522	-0.044020	0.295762	-0.046698	-0.009940	-0.008567	0.044123	1.000000
exng	0.096801	0.141664	-0.394280	0.067616	0.067023	0.025665	-0.070733	-0.378817
caa	0.276326	0.118261	-0.181053	0.101389	0.070511	0.137979	-0.072042	-0.213177
output	-0.225439	-0.280937	0.433798	-0.144931	-0.085239	-0.028046	0.137230	0.421747

```
sns.heatmap(df.corr())
```



<Axes: >




As we can see our variables are not highly correlated to each other

✓ We will do Uni and Bi variate analysis on our Features


```
plt.figure(figsize=(20, 10))
plt.title("Age of Patients")
plt.xlabel("Age")
sns.countplot(x='age',data=df)
```

➡ <Axes: title={'center': 'Age of Patients'}, xlabel='Age', ylabel='count'>




✓ As we can see the Patients are of Age Group 51-67years in majority

```
plt.figure(figsize=(20, 10))
plt.title("Sex of Patients,0=Female and 1=Male")
sns.countplot(x='sex',data=df)
```



```
cp_data = df['cp'].value_counts().reset_index()
cp_data.loc[3, 'index'] = 'asymptomatic'
cp_data.loc[2, 'index'] = 'non-anginal'
cp_data.loc[1, 'index'] = 'Atypical Anigma'
cp_data.loc[0, 'index'] = 'Typical Anigma'
cp_data
```




	cp	count	index
0	0	143	Typical Anigma
1	2	87	Atypical Anigma
2	1	50	non-anginal
3	3	23	asymptomatic

```
plt.figure(figsize=(20, 10))
plt.title("Chest Pain of Patients")
```

```
sns.barplot(x=cp_data['index'],y= cp_data['cp'])
```

```
<Axes: title={'center': 'Chest Pain of Patients'}, xlabel='index', ylabel='cp'>
```



✓ We have seen how the the Chest Pain Category is distributed

```
ecg_data= df['restecg'].value_counts().reset_index()
ecg_data[0,'index']= 'normal'
ecg_data[1,'index']= 'having ST-T wave abnormality'
ecg_data[2,'index']= 'showing probable or definite left ventricular hypertrophy by Estes'
ecg_data
```




	restecg	count	(0, index)	(1, index)	(2, index)
0	1	152	normal	having ST-T wave abnormality	showing probable or definite left ventricular ...
1	0	147	normal	having ST-T wave abnormality	showing probable or definite left ventricular ...

```
ecg_data = df['restecg'].value_counts().reset_index()
ecg_data.columns = ['restecg_type', 'count']
```

```
plt.figure(figsize=(20, 10))
plt.title("ECG data of Patients")
sns.barplot(x=ecg_data['restecg_type'], y=ecg_data['count'])
```



```
<Axes: title={'center': 'ECG data of Patients'}, xlabel='restecg_type',
ylabel='count'>
```

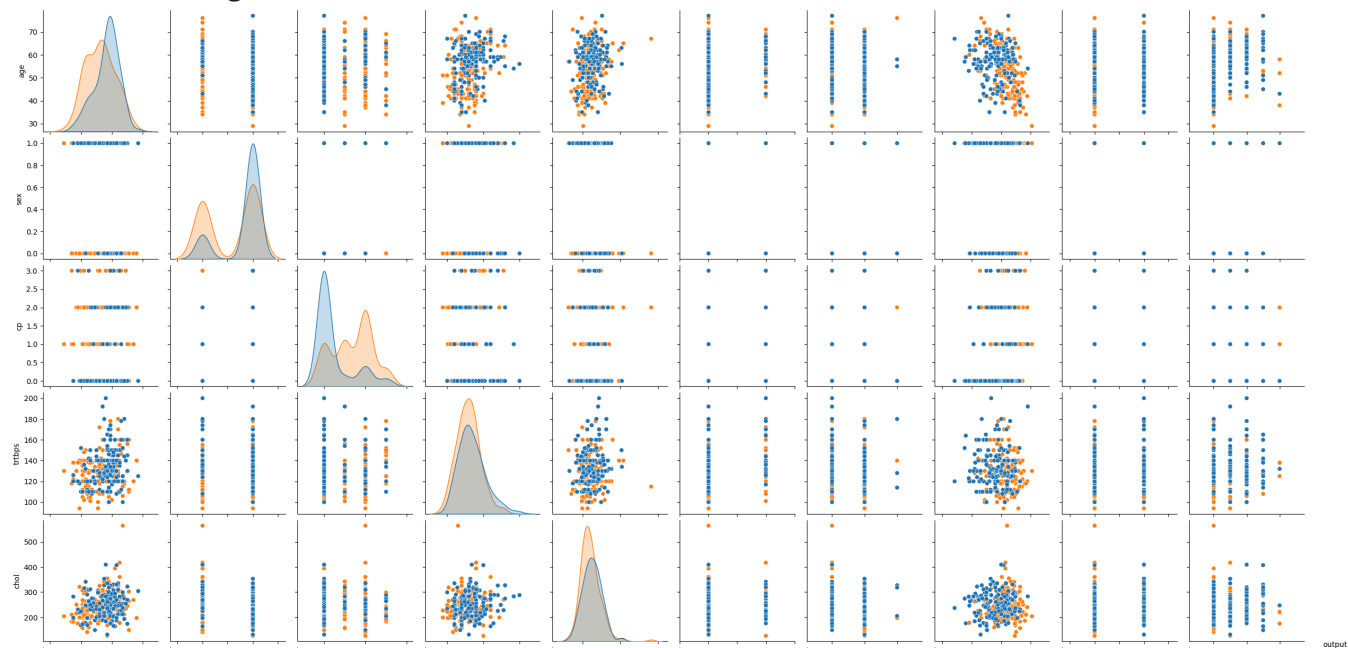


```
sns.pairplot(data=df, hue='output')
```





```
<seaborn.axisgrid.PairGrid at 0x78a437eae610>
```



## ✓ Let us see for our Continuous Variable

```
plt.figure(figsize=(20,10))
plt.subplot(1,2,1)
sns.distplot(df['trtbps'], kde=True, color = 'magenta')
plt.xlabel("Resting Blood Pressure (mmHg)")
plt.subplot(1,2,2)
sns.distplot(df['thalachh'], kde=True, color = 'teal')
plt.xlabel("Maximum Heart Rate Achieved (bpm)")
```

➡ <ipython-input-35-17a8725cb836>:3: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

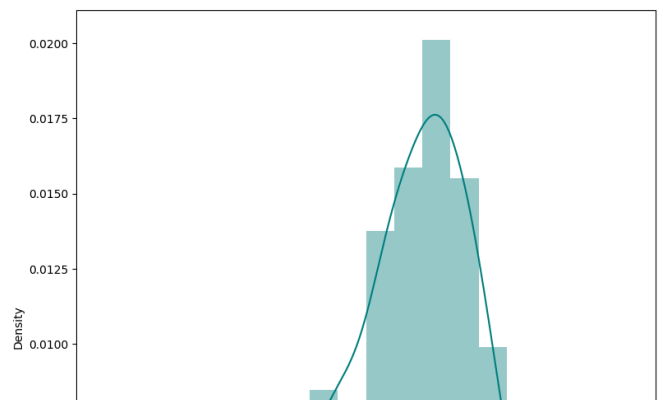

```
sns.distplot(df['trtbps'], kde=True, color = 'magenta')
<ipython-input-35-17a8725cb836>:6: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['thalachh'], kde=True, color = 'teal')
Text(0.5, 0, 'Maximum Heart Rate Achieved (bpm)')
```



```
plt.figure(figsize=(10,10))
sns.distplot(df['chol'], kde=True, color = 'red')
plt.xlabel("Cholestrol")
```


↳ <ipython-input-36-ebe894739d0c>:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['chol'], kde=True, color = 'red')
Text(0.5, 0, 'Cholestrol')
```



- ✓ We have done the Analysis of the data now let's have a look at our data

```
df.head()
```




	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	caa	output
0	63	1	3	145	233	1	0	150	0	0	1
1	37	1	2	130	250	0	1	187	0	0	1
2	41	0	1	130	204	0	0	172	0	0	1
3	56	1	1	120	236	0	1	178	0	0	1
4	57	0	0	120	354	0	1	163	1	0	1

- ✓ Let us do Standardisation

```
from sklearn.preprocessing import StandardScaler
```

```
scale=StandardScaler()
```

```
scale.fit(df)
```



▼ StandardScaler ⓘ ?

StandardScaler()

```
df= scale.transform(df)
```

```
df=pd.DataFrame(df,columns=['age', 'sex', 'cp', 'trtbps', 'chol', 'fbs', 'restecg', 'thalac', 'exng', 'caa', 'output'])
```

```
df.head()
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	output
0	0.952197	0.681005	1.973123	0.763956	-0.256334	2.394438	-1.005832	0.015443	-0.69
1	-1.915313	0.681005	1.002577	-0.092738	0.072199	-0.417635	0.898962	1.633471	-0.69
2	-1.474158	-1.468418	0.032031	-0.092738	-0.816773	-0.417635	-1.005832	0.977514	-0.69
3	0.180175	0.681005	0.032031	-0.663867	-0.198357	-0.417635	0.898962	1.239897	-0.69
4	0.290464	-1.468418	-0.938515	-0.663867	2.082050	-0.417635	0.898962	0.583939	1.44

We can insert this data into our ML Models

✓ We will use the following models for our predictions :

- Logistic Regression
- Decision Tree
- Random Forest
- K Nearest Neighbour
- SVM

Then we will use the ensembling techniques

✓ Let us split our data

```
x= df.iloc[:, :-1]
x
```



	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	
0	0.952197	0.681005	1.973123	0.763956	-0.256334	2.394438	-1.005832	0.015443	-0
1	-1.915313	0.681005	1.002577	-0.092738	0.072199	-0.417635	0.898962	1.633471	-0
2	-1.474158	-1.468418	0.032031	-0.092738	-0.816773	-0.417635	-1.005832	0.977514	-0
3	0.180175	0.681005	0.032031	-0.663867	-0.198357	-0.417635	0.898962	1.239897	-0
4	0.290464	-1.468418	-0.938515	-0.663867	2.082050	-0.417635	0.898962	0.583939	1
...	...	...	...	...	...	...	...	...	...
298	0.290464	-1.468418	-0.938515	0.478391	-0.101730	-0.417635	0.898962	-1.165281	1
299	-1.033002	0.681005	1.973123	-1.234996	0.342756	-0.417635	0.898962	-0.771706	-0
300	1.503641	0.681005	-0.938515	0.706843	-1.029353	2.394438	0.898962	-0.378132	-0
301	0.290464	0.681005	-0.938515	-0.092738	-2.227533	-0.417635	0.898962	-1.515125	1
302	0.290464	-1.468418	0.032031	-0.092738	-0.198357	-0.417635	-1.005832	1.064975	-0

303 rows × 10 columns

```
y= df.iloc[:,-1:]  
y
```



	output
0	0.914529
1	0.914529
2	0.914529
3	0.914529
4	0.914529
...	...
298	-1.093459
299	-1.093459
300	-1.093459
301	-1.093459
302	-1.093459

303 rows × 1 columns

```
from sklearn.model_selection import train_test_split  
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=101)
```

## ▼ Logistic Regression

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.preprocessing import LabelEncoder
```

```
lbl= LabelEncoder()
```

```
encoded_y= lbl.fit_transform(y_train)
```

```
➞ /usr/local/lib/python3.11/dist-packages/sklearn/preprocessing/_label.py:110: Dat  
y = column_or_1d(y, warn=True)
```

```
logreg= LogisticRegression()
```

```
logreg = LogisticRegression()  
logreg.fit(x_train, encoded_y)
```

```
➞ ▼ LogisticRegression ⓘ ?  
LogisticRegression()
```

```
encoded_y
```

```
➞ array([[0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0,  
1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0,  
1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 1,  
0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0,  
1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0,  
0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1,  
0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1,  
1, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1,  
1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1,  
1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1])
```

```
from sklearn.metrics import accuracy_score  
from sklearn.metrics import confusion_matrix
```

```
encoded_ytest= lbl.fit_transform(y_test)
```

```
➞ /usr/local/lib/python3.11/dist-packages/sklearn/preprocessing/_label.py:110: Dat  
y = column_or_1d(y, warn=True)
```

```
Y_pred1 = logreg.predict(x_test)
lr_conf_matrix = confusion_matrix(encoded_ytest, Y_pred1 )
lr_acc_score = accuracy_score(encoded_ytest, Y_pred1)
```

```
lr_conf_matrix
```

```
↵ array([[35,  9],
        [ 4, 43]])
```

```
print(lr_acc_score*100,"%")
```

```
↵ 85.71428571428571 %
```

As we see the Logistic Regression Model have a 85% accuracy

## ✓ Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
```

```
tree= DecisionTreeClassifier()
```

```
tree.fit(x_train,encoded_y)
```

```
↵ ▾ DecisionTreeClassifier ⓘ ?
   DecisionTreeClassifier()
```

```
ypred2=tree.predict(x_test)
```

```
encoded_ytest= lbl.fit_transform(y_test)
```

```
↵ /usr/local/lib/python3.11/dist-packages/sklearn/preprocessing/_label.py:110: Dat
   y = column_or_1d(y, warn=True)
```

```
tree_conf_matrix = confusion_matrix(encoded_ytest,ypred2 )
tree_acc_score = accuracy_score(encoded_ytest, ypred2)
```

```
tree_conf_matrix
```

```
↵ array([[27, 17],
        [10, 37]])
```

```
print(tree_acc_score*100,"%")
```



⇒ 70.32967032967034 %

As we see our Decision Tree Model does not perform well as it gives a score of only 69%

## ✓ Random Forest

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf= RandomForestClassifier()
```

```
rf.fit(x_train,encoded_y)
```

⇒

▼ RandomForestClassifier ⓘ ?

RandomForestClassifier()

```
ypred3 = rf.predict(x_test)
```

```
rf_conf_matrix = confusion_matrix(encoded_ytest,ypred3 )  
rf_acc_score = accuracy_score(encoded_ytest, ypred3)
```

```
rf_conf_matrix
```

⇒

```
array([[31, 13],  
       [ 5, 42]])
```

```
print(rf_acc_score*100,"%")
```

⇒ 80.21978021978022 %

RF also gives us an accuracy of around 80%

## ✓ K Nearest Neighbour

- ✓ We have to select what k we will use for the maximum accuracy

Let's write a function for it


```
from sklearn.neighbors import KNeighborsClassifier
```

```

error_rate= []
for i in range(1,40):
    knn= KNeighborsClassifier(n_neighbors=i)
    knn.fit(x_train,encoded_y)
    pred= knn.predict(x_test)
    error_rate.append(np.mean(pred != encoded_ytest))

plt.figure(figsize=(10,6))
plt.plot(range(1,40),error_rate,color='blue', linestyle='dashed', marker='o',
         markerfacecolor='red', markersize=10)
plt.xlabel('K Vlaue')
plt.ylabel('Error rate')
plt.title('To check the correct value of k')
plt.show()

```



✓ As we see from the graph we should select K= 12 as it gives the best error rate

```

knn= KNeighborsClassifier(n_neighbors=12)
knn.fit(x_train,encoded_y)

```

```
ypred4= knn.predict(x_test)
```

```
knn_conf_matrix = confusion_matrix(encoded_ytest,ypred4 )  
knn_acc_score = accuracy_score(encoded_ytest, ypred4)
```

```
knn_conf_matrix
```

```
⇒ array([[35,  9],  
        [ 5, 42]])
```

```
print(knn_acc_score*100,"%")
```

```
⇒ 84.61538461538461 %
```

As we see KNN gives us an accuracy of around 85% which is good

## ✓ Support Vector Machine(SVM)

```
from sklearn import svm
```

```
svm= svm.SVC()
```

```
svm.fit(x_train,encoded_y)
```

```
⇒ 

▼ SVC ⓘ ?



SVC()


```

```
ypred5= svm.predict(x_test)
```

```
svm_conf_matrix = confusion_matrix(encoded_ytest,ypred5)  
svm_acc_score = accuracy_score(encoded_ytest, ypred5)
```

```
svm_conf_matrix
```

```
⇒ array([[34, 10],  
        [ 8, 39]])
```

```
print(svm_acc_score*100,"%")
```

```
⇒ 80.21978021978022 %
```

We get an accuracy of 80% in SVM

## ✓ Let us see our model accuracy in Table form

```
model_acc= pd.DataFrame({'Model' : ['Logistic Regression','Decision Tree','Random Forest','
```

```
model_acc = model_acc.sort_values(by=['Accuracy'],ascending=False)
```

model\_acc



	Model	Accuracy
0	Logistic Regression	85.714286
3	K Nearest Neighbor	84.615385
2	Random Forest	80.219780
4	SVM	80.219780
1	Decision Tree	70.329670

Let us use one more Techniques known as Adaboost, this is a Boosting technique which uses multiple models for better accuracy.

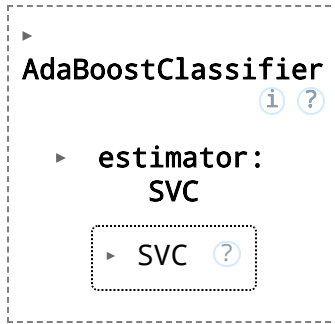
## ✓ Let us first use some random parameters for training the model without Hypertuning.

```
from sklearn.ensemble import AdaBoostClassifier
```

```
adab = AdaBoostClassifier(estimator=svm, n_estimators=100, algorithm='SAMME',  
                           learning_rate=0.01, random_state=0)
```

```
adab.fit(x_train,encoded_y)
```

```
➞ /usr/local/lib/python3.11/dist-packages/sklearn/ensemble/_weight_boosting.py:519
warnings.warn(
```



```
ypred6=adab.predict(x_test)
```

```
adab_conf_matrix = confusion_matrix(encoded_ytest,ypred6)
adab_acc_score = accuracy_score(encoded_ytest, ypred6)
```

```
adab_conf_matrix
```

```
➞ array([[32, 12],
        [ 4, 43]])
```

```
print(adab_acc_score*100,"%")
```

```
➞ 82.41758241758241 %
```

```
adab.score(x_train,encoded_y)
```

```
➞ 0.8160377358490566
```

```
adab.score(x_test,encoded_ytest)
```

```
➞ 0.8241758241758241
```

As we see our model has performed very not bad with just 80% accuracy

We will use Grid Search CV for HyperParameter Tuning

## ✓ Grid Search CV

- ✓ Let us try Grid Search CV for our top 3 performing Algorithms for HyperParameter tuning

```
from sklearn.model_selection import GridSearchCV
```

```
model_acc
```



	Model	Accuracy
0	Logistic Regression	85.714286
3	K Nearest Neighbor	84.615385
2	Random Forest	80.219780
4	SVM	80.219780
1	Decision Tree	70.329670

- ✓ Logistic Regression

```
param_grid= {  
    'solver': ['newton-cg', 'lbfgs', 'liblinear','sag', 'saga'],  
    'penalty' : ['none', 'l1', 'l2', 'elasticnet'],  
    'C' : [100, 10, 1.0, 0.1, 0.01]  
}
```

```
grid1= GridSearchCV(LogisticRegression(),param_grid)
```

```
grid1.fit(x_train,encoded_y)
```

➡ /usr/local/lib/python3.11/dist-packages/sklearn/model\_selection/\_validation.py:5  
325 fits failed out of a total of 500.

The score on these train-test partitions for these parameters will be set to nan  
If these failures are not expected, you can try to debug them by setting error\_s

Below are more details about the failures:

-----  
125 fits failed with the following error:

Traceback (most recent call last):

```
File "/usr/local/lib/python3.11/dist-packages/sklearn/model_selection/_validation.py", line 1382, in _validate_estimator
    estimator.fit(X_train, y_train, **fit_params)
File "/usr/local/lib/python3.11/dist-packages/sklearn/base.py", line 436, in _validate_estimator
    estimator._validate_params()
File "/usr/local/lib/python3.11/dist-packages/sklearn/base.py", line 436, in _validate_estimator
    validate_parameter_constraints(
File "/usr/local/lib/python3.11/dist-packages/sklearn/utils/_param_validation.py", line 100, in _raise_invalid_parameter_error
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'penalty' parameter o
```

-----  
25 fits failed with the following error:

Traceback (most recent call last):

```
File "/usr/local/lib/python3.11/dist-packages/sklearn/model_selection/_validation.py", line 1389, in _validate_estimator
    estimator.fit(X_train, y_train, **fit_params)
File "/usr/local/lib/python3.11/dist-packages/sklearn/base.py", line 436, in _validate_estimator
    estimator._validate_params()
File "/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py", line 1389, in _validate_estimator
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py", line 1389, in _validate_estimator
    raise ValueError(
ValueError: Solver newton-cg supports only 'l2' or None penalties, got l1 penalty
```

-----  
25 fits failed with the following error:

Traceback (most recent call last):

```
File "/usr/local/lib/python3.11/dist-packages/sklearn/model_selection/_validation.py", line 1389, in _validate_estimator
    estimator.fit(X_train, y_train, **fit_params)
File "/usr/local/lib/python3.11/dist-packages/sklearn/base.py", line 436, in _validate_estimator
    estimator._validate_params()
File "/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py", line 1389, in _validate_estimator
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py", line 1389, in _validate_estimator
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or None penalties, got l1 penalty.
```

-----  
25 fits failed with the following error:

Traceback (most recent call last):

```
File "/usr/local/lib/python3.11/dist-packages/sklearn/model_selection/_validation.py", line 1389, in _validate_estimator
    estimator.fit(X_train, y_train, **fit_params)
File "/usr/local/lib/python3.11/dist-packages/sklearn/base.py", line 436, in _validate_estimator
    estimator._validate_params()
File "/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py", line 1389, in _validate_estimator
    solver = _check_solver(self.solver, self.penalty, self.dual)
File "/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py", line 1389, in _validate_estimator
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or None penalties, got l1 penalty.
```





[illegible][illegible]

```
4. let solver = check_solver(self.solver, self.penalty, self.dual)
```

#####

```
logreg1= LogisticRegression(C=0.01,penalty='l2',solver='liblinear')
```

Feedback (most recent call last):

```
LogisticRegression(C=0.01, solver='liblinear')
return fit_method(estimator, *args, **kwargs)
```

```
raise ValueError("l1_ratio must be specified when penalty is elasticnet.")
```

```
logreg_pred_acc_score = accuracy_score(encoded_ytest, logreg_pred)
```

logreg\_pred\_conf\_matrix

```
>> array([[nan, nan], [nan, nan]])
```

0 70046700

⇒	80.38316732868131 %nan	nan	nan	nan	nan
	nan nan	nan	nan	nan	nan
	nan 0.77364341	nan 0.76877076	0.78781838	0.78781838	

nan	nan	nan	nan	nan	nan
-----	-----	-----	-----	-----	-----

0 78216722 0 78216722 0 70756268 0 78216722 0 78216722 nan

```

KINN
warnings.warn(

```

► GridSearchCV

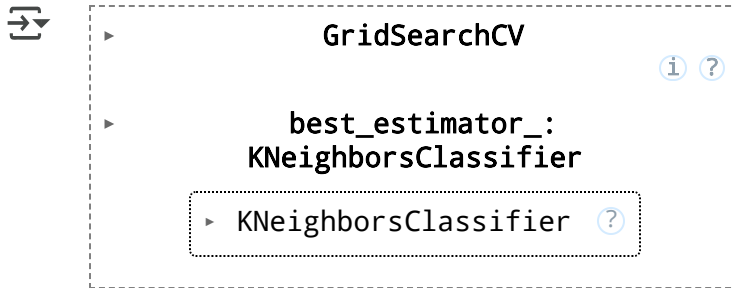
```
weights = ['uniform', 'distance']
```

```
metric = ['euclidean', 'manhattan', 'minkowski']
```

For \_\_\_\_\_

```
grid_search = GridSearchCV(estimator=knn, param_grid=grid, n_jobs=-1, cv=cv, scoring='accuracy')
```

```
grid_search.fit(x_train,encoded_y)
```



```
grid_search.best_params_
```

```
{'metric': 'manhattan', 'n_neighbors': 11, 'weights': 'distance'}
```

## ✓ Let's apply

```
knn= KNeighborsClassifier(n_neighbors=12,metric='manhattan',weights='distance')
knn.fit(x_train,encoded_y)
knn_pred= knn.predict(x_test)
```

```
knn_pred_conf_matrix = confusion_matrix(encoded_ytest,knn_pred)
knn_pred_acc_score = accuracy_score(encoded_ytest, knn_pred)
```

```
knn_pred_conf_matrix
```

```
array([[33, 11],
       [ 5, 42]])
```

```
print(knn_pred_acc_score*100,"%")
```

```
82.41758241758241 %
```


We have an Accuracy of 82.5%

## ✓ SVM

```
kernel = ['poly', 'rbf', 'sigmoid']
C = [50, 10, 1.0, 0.1, 0.01]
gamma = ['scale']
```

```
grid = dict(kernel=kernel,C=C,gamma=gamma)
cv = RepeatedStratifiedKFold(n_splits=10, n_repeats=3, random_state=1)
grid_search = GridSearchCV(estimator=svm, param_grid=grid, n_jobs=-1, cv=cv, scoring='accuracy')
```

```
grid_search.fit(x_train,encoded_y)
```



```
grid_search.best_params_
```

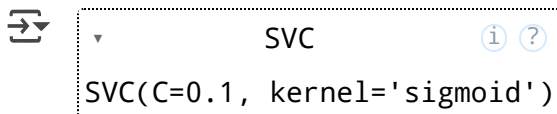
```
{'C': 0.1, 'gamma': 'scale', 'kernel': 'sigmoid'}
```

✓ Let us apply these

```
from sklearn.svm import SVC
```

```
svc= SVC(C= 0.1, gamma= 'scale',kernel= 'sigmoid')
```

```
svc.fit(x_train,encoded_y)
```



```
svm_pred= svc.predict(x_test)
```

```
svm_pred_conf_matrix = confusion_matrix(encoded_ytest,svm_pred)
svm_pred_acc_score = accuracy_score(encoded_ytest, svm_pred)
```

```
svm_pred_conf_matrix
```

```
array([[32, 12],
       [ 5, 42]])
```

```
print(svm_pred_acc_score*100,"%")
```

```
81.31868131868131 %
```

Accuracy is 81%

## ✓ Final Verdict

✓ After comparing all the models the best performing model is :

## Logistic Regression with no Hyperparameter tuning

```
logreg= LogisticRegression()  
logreg = LogisticRegression()  
logreg.fit(x_train, encoded_y)
```



```
▼ LogisticRegression ⓘ ?  
LogisticRegression()
```

Y\_pred1



```
array([0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1,  
       0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1,  
       1, 0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0,  
       1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1,  
       0, 1, 0])
```

lr\_conf\_matrix



```
array([[35,  9],  
       [ 4, 43]])
```

```
print(lr_acc_score*100,"%")
```



```
85.71428571428571 %
```

✓ Let us build a proper confusion matrix for our model

```
# Confusion Matrix of Model enlarged  
options = ["Disease", 'No Disease']
```

```
fig, ax = plt.subplots()  
im = ax.imshow(lr_conf_matrix, cmap= 'Set3', interpolation='nearest')
```

```
# We want to show all ticks...  
ax.set_xticks(np.arange(len(options)))  
ax.set_yticks(np.arange(len(options)))  
# ... and label them with the respective list entries
```

```

ax.set_xticklabels(options)
ax.set_yticklabels(options)

# Rotate the tick labels and set their alignment.
plt.setp(ax.get_xticklabels(), rotation=45, ha="right",
         rotation_mode="anchor")

# Loop over data dimensions and create text annotations.
for i in range(len(options)):
    for j in range(len(options)):
        text = ax.text(j, i, lr_conf_matrix[i, j],
                       ha="center", va="center", color="black")

ax.set_title("Confusion Matrix of Logistic Regression Model")
fig.tight_layout()
plt.xlabel('Model Prediction')
plt.ylabel('Actual Result')
plt.show()
print("ACCURACY of our model is ",lr_acc_score*100,"%")

```



- We have successfully made our model which predicts whether a
- ✓ person is having a risk of Heart Disease or not with 85.7% accuracy

```
import pickle
```

```
pickle.dump(logreg,open('heart.pkl','wb'))
```

## Using Auto ML

EVAL ML :



EvalML is an open-source AutoML library written in python that automates a large part of the machine learning process and we can easily evaluate which machine learning pipeline works better for the given set of data.

- ✓ Installing Eval ML

```
!pip install evalml
```

```
Requirement already satisfied: evalml in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: numpy>=1.22.0 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: pandas<2.1.0,>=1.5.0 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: scipy>=1.5.0 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: scikit-learn>=1.3.2 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: scikit-optimize>=0.9.0 in /usr/local/lib/python3.11/dist-packages
Requirement already satisfied: pyzmq>=20.0.0 in /usr/local/lib/python3.11/dist-packages
```

```


Requirement already satisfied: colorama>=0.4.4 in /usr/local/lib/python3.11/dis
Requirement already satisfied: cloudpickle>=1.5.0 in /usr/local/lib/python3.11/
Requirement already satisfied: click>=8.0.0 in /usr/local/lib/python3.11/dist-p
Requirement already satisfied: shap>=0.45.0 in /usr/local/lib/python3.11/dist-p
Requirement already satisfied: statsmodels>=0.12.2 in /usr/local/lib/python3.11
Requirement already satisfied: texttable>=1.6.2 in /usr/local/lib/python3.11/di
Requirement already satisfied: woodwork>=0.22.0 in /usr/local/lib/python3.11/di
Requirement already satisfied: dask!=2022.10.1,>=2022.2.0 in /usr/local/lib/pyt
Requirement already satisfied: distributed!=2022.10.1,>=2022.2.0 in /usr/local/
Requirement already satisfied: featuretools>=1.16.0 in /usr/local/lib/python3.1
Requirement already satisfied: nlp-primitives>=2.9.0 in /usr/local/lib/python3.
Requirement already satisfied: networkx>=2.7 in /usr/local/lib/python3.11/dist-
Requirement already satisfied: plotly>=5.0.0 in /usr/local/lib/python3.11/dist-
Requirement already satisfied: kaleido>=0.2.0 in /usr/local/lib/python3.11/dist
Requirement already satisfied: ipywidgets>=7.5 in /usr/local/lib/python3.11/dis
Requirement already satisfied: xgboost>=1.7.0.post0 in /usr/local/lib/python3.1
Requirement already satisfied: catboost>=1.1.1 in /usr/local/lib/python3.11/dis
Requirement already satisfied: lightgbm>=4.0.0 in /usr/local/lib/python3.11/dis
Requirement already satisfied: matplotlib>=3.3.3 in /usr/local/lib/python3.11/c
Requirement already satisfied: seaborn>=0.11.1 in /usr/local/lib/python3.11/dis
Requirement already satisfied: category-encoders<=2.5.1.post0,>=2.2.2 in /usr/l
Requirement already satisfied: imbalanced-learn>=0.11.0 in /usr/local/lib/pytho
Requirement already satisfied: pmdarima>=1.8.5 in /usr/local/lib/python3.11/dis
Requirement already satisfied: sktime<0.29.0,>=0.21.0 in /usr/local/lib/python3
Requirement already satisfied: lime>=0.2.0.1 in /usr/local/lib/python3.11/dist-
Requirement already satisfied: toml>=2.0.1 in /usr/local/lib/python3.11/dist-p
Requirement already satisfied: packaging>=23.0 in /usr/local/lib/python3.11/dis
Requirement already satisfied: black>=22.3.0 in /usr/local/lib/python3.11/dist-
Requirement already satisfied: holidays>=0.13 in /usr/local/lib/python3.11/dist
Requirement already satisfied: graphviz>=0.13 in /usr/local/lib/python3.11/dist
Requirement already satisfied: mypy-extensions>=0.4.3 in /usr/local/lib/python3
Requirement already satisfied: pathspec>=0.9.0 in /usr/local/lib/python3.11/dis
Requirement already satisfied: platformdirs>=2 in /usr/local/lib/python3.11/dis
Requirement already satisfied: ipython>=7.8.0 in /usr/local/lib/python3.11/dist
Requirement already satisfied: tokenize-rt>=3.2.0 in /usr/local/lib/python3.11/
Requirement already satisfied: six in /usr/local/lib/python3.11/dist-packages (
Requirement already satisfied: patsy>=0.5.1 in /usr/local/lib/python3.11/dist-p
Requirement already satisfied: fsspec>=2021.09.0 in /usr/local/lib/python3.11/c
Requirement already satisfied: partd>=1.4.0 in /usr/local/lib/python3.11/dist-p
Requirement already satisfied: pyyaml>=5.3.1 in /usr/local/lib/python3.11/dist-
Requirement already satisfied: toolz>=0.10.0 in /usr/local/lib/python3.11/dist-
Requirement already satisfied: importlib_metadata>=4.13.0 in /usr/local/lib/pyt
Requirement already satisfied: jinja2>=2.10.3 in /usr/local/lib/python3.11/dist
Requirement already satisfied: lock>=1.0.0 in /usr/local/lib/python3.11/dist-
Requirement already satisfied: msgpack>=1.0.2 in /usr/local/lib/python3.11/dist
Requirement already satisfied: psutil>=5.8.0 in /usr/local/lib/python3.11/dist-
Requirement already satisfied: sortedcontainers>=2.0.5 in /usr/local/lib/pythor
Requirement already satisfied: tblib>=1.6.0 in /usr/local/lib/python3.11/dist-p
Requirement already satisfied: tornado>=6.2.0 in /usr/local/lib/python3.11/dist
Requirement already satisfied: urllib3>=1.26.5 in /usr/local/lib/python3.11/dis

```

✓ *Let us load our DataSet.*


```
import pandas as pd
df= pd.read_csv("/content/heart.csv")
```

```
df.head()
```



	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thal
<b>0</b>	63	1	3	145	233	1	0	150	0	2.3	0	0	
<b>1</b>	37	1	2	130	250	0	1	187	0	3.5	0	0	
<b>2</b>	41	0	1	130	204	0	0	172	0	1.4	2	0	
<b>3</b>	56	1	1	120	236	0	1	178	0	0.8	2	0	
<b>4</b>	57	0	0	120	354	0	1	163	1	0.6	2	0	

```
x= df.iloc[:, :-1]
x
```



	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	th
<b>0</b>	63	1	3	145	233	1	0	150	0	2.3	0	0	
<b>1</b>	37	1	2	130	250	0	1	187	0	3.5	0	0	
<b>2</b>	41	0	1	130	204	0	0	172	0	1.4	2	0	
<b>3</b>	56	1	1	120	236	0	1	178	0	0.8	2	0	
<b>4</b>	57	0	0	120	354	0	1	163	1	0.6	2	0	
...	...	...	...	...	...	...	...	...	...	...	...	...	...
<b>298</b>	57	0	0	140	241	0	1	123	1	0.2	1	0	
<b>299</b>	45	1	3	110	264	0	1	132	0	1.2	1	0	
<b>300</b>	68	1	0	144	193	1	1	141	0	3.4	1	2	
<b>301</b>	57	1	0	130	131	0	1	115	1	1.2	1	1	
<b>302</b>	57	0	1	130	236	0	0	174	0	0.0	1	1	

303 rows × 13 columns

```
import pandas as pd
from sklearn.preprocessing import LabelEncoder
```

```
# Instantiate the LabelEncoder
lbl = LabelEncoder()
```

```
y= df.iloc[:, -1:]
y= lbl.fit_transform(y) # Use lbl instead of lb
y
```



```

➡ /usr/local/lib/python3.11/dist-packages/sklearn/preprocessing/_label.py:114: Dat
y = column_or_1d(y, warn=True)
array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])

```

## ✓ Importing Eval ML Library

```
import evalml
```

Eval ML Library will do all the pre processing techniques for us and split the data for us

```
X_train, X_test, y_train, y_test = evalml.preprocessing.split_data(x, y, problem_type='bina
```

There are different problem type parameters in Eval ML, we have a Binary type problem here, that's why we are using Binary as a input

## Running the Auto ML to select best Algorithm

```
from evalml automl import AutoMLSearch
automl = AutoMLSearch(X_train=X_train, y_train=y_train, problem_type='binary')
automl.search()
```

```


➡ {1: {'Random Forest Classifier w/ Label Encoder + Imputer + RF Classifier Select
From Model': 2.6637513637542725,
      'Total time of batch': 2.7914462089538574},
   2: {'LightGBM Classifier w/ Label Encoder + Imputer + Select Columns
Transformer': 1.298210859298706,
      'Extra Trees Classifier w/ Label Encoder + Imputer + Select Columns
Transformer': 1.9191703796386719,
      'Elastic Net Classifier w/ Label Encoder + Imputer + Standard Scaler + Select
Columns Transformer': 2.190389394760132,
      'XGBoost Classifier w/ Label Encoder + Imputer + Select Columns Transformer':
2.1611223220825195,
      'Logistic Regression Classifier w/ Label Encoder + Imputer + Standard Scaler +

```

```
Select Columns Transformer': 3.439563274383545,
'Total time of batch': 11.627403020858765}}
```


As we see from the above output the Auto ML Classifier has given us the best fit Algorithm which is Extra Trees Classifier with Imputer We can also compare the rest of the models

```
automl.rankings
```



	id	pipeline_name	search_order	ranking_score	mean_cv_score	standard_deviat:
0	3	Extra Trees Classifier w/ Label Encoder + Impu...	3	0.413358	0.413358	
1	2	LightGBM Classifier w/ Label Encoder + Imputer...	2	0.462099	0.462099	
2	1	Random Forest Classifier w/ Label Encoder + Im...	1	0.466918	0.466918	
3	6	Logistic Regression Classifier w/ Label Encode...	6	0.469250	0.469250	
4	4	Elastic Net Classifier w/ Label Encoder + Impu...	4	0.470037	0.470037	
5	5	XGBoost Classifier w/ Label Encoder + Imputer ...	5	0.499149	0.499149	
6	0	Mode Baseline Binary Classification Pipeline	0	16.382805	16.382805	

```
automl.best_pipeline
```



```
pipeline = BinaryClassificationPipeline(component_graph={'Label Encoder':
['Label Encoder', 'X', 'y'], 'Imputer': ['Imputer', 'X', 'Label Encoder.y'],
'Select Columns Transformer': ['Select Columns Transformer', 'Imputer.x', 'Label
Encoder.y'], 'Extra Trees Classifier': ['Extra Trees Classifier', 'Select
Columns Transformer.x', 'Label Encoder.y']}}, parameters={'Label Encoder':
```

```
{'positive_label': None}, 'Imputer':{'categorical_impute_strategy':
'most_frequent', 'numeric_impute_strategy': 'mean', 'boolean_impute_strategy':
'most_frequent', 'categorical_fill_value': None, 'numeric_fill_value': None,
'boolean_fill_value': None}, 'Select Columns Transformer':{'columns': ['age',
'cp', 'thalachh', 'exng', 'oldpeak', 'caa', 'thall']}}, 'Extra Trees Classifier':
{'n_estimators': 100, 'max_features': 'sqrt', 'max_depth': 6,
'min_samples_split': 2, 'min_weight_fraction_leaf': 0.0, 'n_jobs': -1}},
random_seed=0)
```

```
best_pipeline=automl.best_pipeline
```

We can have a Detailed description of our Best Selected Model

```
automl.describe_pipeline(automl.rankings.iloc[0]["id"])
```



```
*****
INFO:evalml.pipelines.pipeline_base.describe:
*****
* Extra Trees Classifier w/ Label Encoder + Imputer + Select Columns Transform
INFO:evalml.pipelines.pipeline_base.describe:* Extra Trees Classifier w/ Label
*****
INFO:evalml.pipelines.pipeline_base.describe:*****

INFO:evalml.pipelines.pipeline_base.describe:
Problem Type: binary
INFO:evalml.pipelines.pipeline_base.describe:Problem Type: binary
Model Family: Extra Trees
INFO:evalml.pipelines.pipeline_base.describe:Model Family: Extra Trees

INFO:evalml.pipelines.pipeline_base.describe:
Pipeline Steps
INFO:evalml.pipelines.pipeline_base.describe:Pipeline Steps
=====
INFO:evalml.pipelines.pipeline_base.describe:=====
1. Label Encoder
INFO:evalml.pipelines.component_graph.describe:1. Label Encoder
    * positive_label : None
INFO:evalml.pipelines.components.component_base.describe:    * positive_label
2. Imputer
INFO:evalml.pipelines.component_graph.describe:2. Imputer
    * categorical_impute_strategy : most_frequent
INFO:evalml.pipelines.components.component_base.describe:    * categorical_
    * numeric_impute_strategy : mean
INFO:evalml.pipelines.components.component_base.describe:    * numeric_impu
    * boolean_impute_strategy : most_frequent
INFO:evalml.pipelines.components.component_base.describe:    * boolean_impu
    * categorical_fill_value : None
INFO:evalml.pipelines.components.component_base.describe:    * categorical_
    * numeric_fill_value : None
INFO:evalml.pipelines.components.component_base.describe:    * numeric_fill
    * boolean_fill_value : None
INFO:evalml.pipelines.components.component_base.describe:    * boolean_fill
```

### 3. Select Columns Transformer

```
INFO:evalml.pipelines.component_graph.describe:3. Select Columns Transformer
      * columns : ['age', 'cp', 'thalachh', 'exng', 'oldpeak', 'caa', 'thall']
INFO:evalml.pipelines.components.component_base.describe:      * columns : ['
```

### 4. Extra Trees Classifier

```
INFO:evalml.pipelines.component_graph.describe:4. Extra Trees Classifier
```

```
      * n_estimators : 100
INFO:evalml.pipelines.components.component_base.describe:      * n_estimators
      * max_features : sqrt
INFO:evalml.pipelines.components.component_base.describe:      * max_features
      * max_depth : 6
INFO:evalml.pipelines.components.component_base.describe:      * max_depth :
      * min_samples_split : 2
INFO:evalml.pipelines.components.component_base.describe:      * min_samples_
      * min_weight_fraction_leaf : 0.0
INFO:evalml.pipelines.components.component_base.describe:      * min_weight_f
      * n_jobs : -1
INFO:evalml.pipelines.components.component_base.describe:      * n_jobs : -1
```

```
best_pipeline.score(X_test, y_test, objectives=["auc", "f1", "Precision", "Recall"])
```

```
➡ OrderedDict([('AUC', 0.8701298701298702),
               ('F1', 0.78125),
               ('Precision', 0.8064516129032258),
               ('Recall', 0.7575757575757576)])
```

Now if we want to build our Model for a specific objective we can do that

```
automl_auc = AutoMLSearch(X_train=X_train, y_train=y_train,
                          problem_type='binary',
                          objective='auc',
                          additional_objectives=['f1', 'precision'],
                          max_batches=1,
                          optimize_thresholds=True)
```

```
automl_auc.search()
```

```
➡ {1: {'Random Forest Classifier w/ Label Encoder + Imputer + RF Classifier Select
      From Model': 4.003447532653809,
      'Total time of batch': 4.1510655879974365}}
```

```
automl_auc.rankings
```



```
id pipeline_name search_order ranking_score mean_cv_score standard_deviation
```

Random Forest

Classifier w/

```
automl_auc.describe_pipeline(automl_auc.rankings.iloc[0]["id"])
```



```
*****Mode Baseline*****
INFO:evalml.pipelines.pipeline_base.describe:0.500000      0.500000
*****Binary Classification Pipeline*****
* Random Forest Classifier w/ Label Encoder + Imputer + RF Classifier Select Fr
INFO:evalml.pipelines.pipeline_base.describe:* Random Forest Classifier w/ Labe
*****
INFO:evalml.pipelines.pipeline_base.describe:*****

INFO:evalml.pipelines.pipeline_base.describe:
Problem Type: binary
INFO:evalml.pipelines.pipeline_base.describe:Problem Type: binary
Model Family: Random Forest
INFO:evalml.pipelines.pipeline_base.describe:Model Family: Random Forest

INFO:evalml.pipelines.pipeline_base.describe:
Pipeline Steps
INFO:evalml.pipelines.pipeline_base.describe:Pipeline Steps
=====
INFO:evalml.pipelines.pipeline_base.describe:=====
1. Label Encoder
INFO:evalml.pipelines.component_graph.describe:1. Label Encoder
    * positive_label : None
INFO:evalml.pipelines.components.component_base.describe:      * positive_label

2. Imputer
INFO:evalml.pipelines.component_graph.describe:2. Imputer
    * categorical_impute_strategy : most_frequent
INFO:evalml.pipelines.components.component_base.describe:      * categorical_impute_strategy : most_frequent
    * numeric_impute_strategy : mean
INFO:evalml.pipelines.components.component_base.describe:      * numeric_impute_strategy : mean
    * boolean_impute_strategy : most_frequent
INFO:evalml.pipelines.components.component_base.describe:      * boolean_impute_strategy : most_frequent
    * categorical_fill_value : None
INFO:evalml.pipelines.components.component_base.describe:      * categorical_fill_value : None
    * numeric_fill_value : None
```